Lecture6_note_data_wrangling

Shukyee Chan

2023-10-12

Outline

- Reshape (long <-> wide) with `pivot_longer` and `pivot_wider`
- Stack tables by row or by column with `bind_rows` and `bind_cols` (or, alternatively, `cbind` and `rbind`)
- Merge two tables with `inner_join`, `full_join`, `left_join`, `right_join`, `semi_join`, and `anti_join`
- Save your outputs

```
library(tidyverse)
```

```
d <- read_csv("_DataPublic_/vdem/1984_2022/vdem_1984_2022_external.csv")
d |> print(n = 3)
```

```
## # A tibble: 6,789 x 211
##   country_name countr~1 count~2  year historic~3 project histo~4 histn~5 codin~6
##   <chr>        <chr>      <dbl> <dbl> <date>       <dbl>   <dbl> <chr>     <dbl>
## 1 Mexico       MEX            3  1984 1984-12-31       0       1 United~    1789
## 2 Mexico       MEX            3  1985 1985-12-31       0       1 United~    1789
## 3 Mexico       MEX            3  1986 1986-12-31       0       1 United~    1789
## # ... with 6,786 more rows, 202 more variables: codingend <dbl>,
## #   codingstart_contemp <dbl>, codingend_contemp <dbl>, codingstart_hist <dbl>,
## #   codingend_hist <dbl>, gapstart1 <dbl>, gapstart2 <dbl>, gapstart3 <dbl>,
## #   gapend1 <dbl>, gapend2 <dbl>, gapend3 <dbl>, gap_index <dbl>,
## #   COWcode <dbl>, e_v2x_api_3C <dbl>, e_v2x_api_4C <dbl>, e_v2x_api_5C <dbl>,
## #   e_v2x_civlib_3C <dbl>, e_v2x_civlib_4C <dbl>, e_v2x_civlib_5C <dbl>,
## #   e_v2x_clphy_3C <dbl>, e_v2x_clphy_4C <dbl>, e_v2x_clphy_5C <dbl>, ...
```

```
d_gdp <- d |>
  select(country_text_id, year, e_gdp, e_gdppc) |>
  rename("gdp" = "e_gdp", "gdppc" = "e_gdppc")

d_gdp |> print(n = 3)
```

```
## # A tibble: 6,789 x 4
##   country_text_id  year    gdp gdppc
##   <chr>           <dbl>  <dbl> <dbl>
## 1 MEX              1984 93563.  11.7
## 2 MEX              1985 94259.  11.5
## 3 MEX              1986 92750.  11.1
## # ... with 6,786 more rows
```

1. Reshape a Table

Wide to Long: `pivot_longer`

```
d_gdp_long <- d_gdp |>
  pivot_longer(cols = c("gdp", "gdppc"),
               names_to = "variable", values_to = "value")

d_gdp_long |> print(n = 4)
```

```
## # A tibble: 13,578 x 4
##   country_text_id  year variable   value
##   <chr>           <dbl> <chr>      <dbl>
## 1 MEX              1984 gdp      93563.
## 2 MEX              1984 gdppc       11.7
## 3 MEX              1985 gdp      94259.
## 4 MEX              1985 gdppc       11.5
## # ... with 13,574 more rows
```

Long to Wide: `pivot_wider`

```
# Reverse the above pivot_long operation
d_gdp_wide_1 <- d_gdp_long |>
  pivot_wider(names_from = "variable", values_from = "value")

d_gdp_wide_1 |> print(n = 4)
```

```
## # A tibble: 6,789 x 4
##   country_text_id  year    gdp gdppc
##   <chr>           <dbl>  <dbl> <dbl>
## 1 MEX              1984 93563.  11.7
## 2 MEX              1985 94259.  11.5
## 3 MEX              1986 92750.  11.1
## 4 MEX              1987 93220.  10.9
## # ... with 6,785 more rows
```

```
# Make year the column variable
d_gdp_wide_2 <- d_gdp_long |>
  pivot_wider(names_from = "year", values_from = "value")

d_gdp_wide_2 |> print(n = 2)
```

```
## # A tibble: 362 x 41
##   count~1 varia~2 `1984` `1985` `1986` `1987` `1988` `1989` `1990` `1991` `1992`
##   <chr>   <chr>    <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1 MEX     gdp     9.36e4 9.43e4 9.28e4 9.32e4 9.47e4 9.81e4 1.03e5 1.07e5 1.12e5
## 2 MEX     gdppc   1.17e1 1.15e1 1.11e1 1.09e1 1.08e1 1.10e1 1.14e1 1.16e1 1.19e1
## # ... with 360 more rows, 30 more variables: `1993` <dbl>, `1994` <dbl>,
## #   `1995` <dbl>, `1996` <dbl>, `1997` <dbl>, `1998` <dbl>, `1999` <dbl>,
```

```
## #   `2000` <dbl>, `2001` <dbl>, `2002` <dbl>, `2003` <dbl>, `2004` <dbl>,
## #   `2005` <dbl>, `2006` <dbl>, `2007` <dbl>, `2008` <dbl>, `2009` <dbl>,
## #   `2010` <dbl>, `2011` <dbl>, `2012` <dbl>, `2013` <dbl>, `2014` <dbl>,
## #   `2015` <dbl>, `2016` <dbl>, `2017` <dbl>, `2018` <dbl>, `2019` <dbl>,
## #   `2020` <dbl>, `2021` <dbl>, `2022` <dbl>, and abbreviated variable ...
```

```
#  Make country_text_id the column variable
d_gdp_wide_3 <- d_gdp_long |>
  pivot_wider(names_from = "country_text_id", values_from = "value")

d_gdp_wide_3 |> print(n = 2)
```

```
## # A tibble: 78 x 183
##    year variable     MEX    SUR    SWE    CHE    GHA    ZAF    JPN    MMR    RUS
##   <dbl> <chr>      <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>  <dbl>
## 1  1984 gdp      93563.  286.   2.35e4 2.31e4 3.02e3 3.15e4 2.87e5 4.18e3 3.49e5
## 2  1984 gdppc       11.7   7.43 2.66e1 3.32e1 2.20e0 9.03e0 2.26e1 1.10e0 1.65e1
## # ... with 76 more rows, and 172 more variables: ALB <dbl>, EGY <dbl>,
## #   YEM <dbl>, COL <dbl>, POL <dbl>, BRA <dbl>, USA <dbl>, PRT <dbl>,
## #   SLV <dbl>, YMD <dbl>, BGD <dbl>, BOL <dbl>, HTI <dbl>, HND <dbl>,
## #   MLI <dbl>, PAK <dbl>, PER <dbl>, SEN <dbl>, SSD <dbl>, SDN <dbl>,
## #   VNM <dbl>, AFG <dbl>, ARG <dbl>, ETH <dbl>, IND <dbl>, KEN <dbl>,
## #   PRK <dbl>, KOR <dbl>, XKX <dbl>, LBN <dbl>, NGA <dbl>, PHL <dbl>,
## #   TZA <dbl>, TWN <dbl>, THA <dbl>, UGA <dbl>, VEN <dbl>, BEN <dbl>, ...
```

2. Stack Tables

```
# New data (for stack data vertically)
d_gdp_1945 <-
  read_csv("_DataPublic_/vdem/1945_1983/vdem_1945_1983_external.csv") |>
  select(country_text_id, year, e_gdp, e_gdppc) |>
  rename("gdp" = "e_gdp", "gdppc" = "e_gdppc")

d_gdp_1906 <-
  read_csv("_DataPublic_/vdem/1906_1944/vdem_1906_1944_external.csv") |>
  select(country_text_id, year, e_gdp, e_gdppc) |>
  rename("gdp" = "e_gdp", "gdppc" = "e_gdppc")

d_gdp_1945 |> print(n = 2)
```

```
## # A tibble: 6,082 x 4
##   country_text_id  year   gdp gdppc
##   <chr>           <dbl> <dbl> <dbl>
## 1 MEX              1945 7827.  3.08
## 2 MEX              1946 8331.  3.17
## # ... with 6,080 more rows
```

```
# New data (for stack data horizontally)
d_edu <- d |>
  select(e_peaveduc, e_peedgini) |>
```

```
  rename("edu_15" = "e_peaveduc", "edu_gini" = "e_peedgini")

d_fh <- d |>
  select(starts_with("e_fh")) |>
  rename("fh_CivilLiberty" = "e_fh_cl", "fh_PoliticalRight" = "e_fh_pr",
         "fh_RuleOfLaw" = "e_fh_rol", "fh_Status" = "e_fh_status")

d_fh |> print(n = 2)
```

```
## # A tibble: 6,789 x 4
##   fh_CivilLiberty fh_PoliticalRight fh_RuleOfLaw fh_Status
##             <dbl>             <dbl>        <dbl>     <dbl>
## 1               4                 3           NA         2
## 2               4                 4           NA         2
## # ... with 6,787 more rows
```

bind_rows

```
d_gdp_1945_2022 <- bind_rows(d_gdp, d_gdp_1945)
d_gdp_1945_2022 |> print(n = 3)
```

```
## # A tibble: 12,871 x 4
##   country_text_id  year    gdp gdppc
##   <chr>           <dbl>  <dbl> <dbl>
## 1 MEX              1984 93563.  11.7
## 2 MEX              1985 94259.  11.5
## 3 MEX              1986 92750.  11.1
## # ... with 12,868 more rows
```

```
unique(d_gdp_1945_2022$year) |> sort()
```

```
##  [1] 1945 1946 1947 1948 1949 1950 1951 1952 1953 1954 1955 1956 1957 1958 1959
## [16] 1960 1961 1962 1963 1964 1965 1966 1967 1968 1969 1970 1971 1972 1973 1974
## [31] 1975 1976 1977 1978 1979 1980 1981 1982 1983 1984 1985 1986 1987 1988 1989
## [46] 1990 1991 1992 1993 1994 1995 1996 1997 1998 1999 2000 2001 2002 2003 2004
## [61] 2005 2006 2007 2008 2009 2010 2011 2012 2013 2014 2015 2016 2017 2018 2019
## [76] 2020 2021 2022
```

```
d_gdp_1945_2022_ue_rows <- bind_rows(
 d_gdp |> select(-gdppc),
 d_gdp_1945 |> select(-gdp)
)
```

```
d_gdp_1906_2022 <- bind_rows(d_gdp, d_gdp_1945, d_gdp_1906)
d_gdp_1906_2022 |> print(n = 3)
```

```
## # A tibble: 18,559 x 4
##   country_text_id  year    gdp gdppc
##   <chr>           <dbl>  <dbl> <dbl>
```

```
## 1 MEX              1984 93563.  11.7
## 2 MEX              1985 94259.  11.5
## 3 MEX              1986 92750.  11.1
## # ... with 18,556 more rows
```

```
unique(d_gdp_1906_2022$year) |> sort()
```

```
##   [1] 1906 1907 1908 1909 1910 1911 1912 1913 1914 1915 1916 1917 1918 1919 1920
##  [16] 1921 1922 1923 1924 1925 1926 1927 1928 1929 1930 1931 1932 1933 1934 1935
##  [31] 1936 1937 1938 1939 1940 1941 1942 1943 1944 1945 1946 1947 1948 1949 1950
##  [46] 1951 1952 1953 1954 1955 1956 1957 1958 1959 1960 1961 1962 1963 1964 1965
##  [61] 1966 1967 1968 1969 1970 1971 1972 1973 1974 1975 1976 1977 1978 1979 1980
##  [76] 1981 1982 1983 1984 1985 1986 1987 1988 1989 1990 1991 1992 1993 1994 1995
##  [91] 1996 1997 1998 1999 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010
## [106] 2011 2012 2013 2014 2015 2016 2017 2018 2019 2020 2021 2022
```

bind_cols

```
d_gdp_edu_fh <- bind_cols(d_gdp, d_edu, d_fh)
d_gdp_edu_fh |> print(n = 3)
```

```
## # A tibble: 6,789 x 10
##   country_te~1  year    gdp gdppc edu_15 edu_g~2 fh_Ci~3 fh_Po~4 fh_Ru~5 fh_St~6
##   <chr>        <dbl>  <dbl> <dbl>  <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 MEX           1984 93563.  11.7   6.08    32.7       4       3      NA       2
## 2 MEX           1985 94259.  11.5   6.22    32.4       4       4      NA       2
## 3 MEX           1986 92750.  11.1   6.36    31.9       4       4      NA       2
## # ... with 6,786 more rows, and abbreviated variable names 1: country_text_id,
## #   2: edu_gini, 3: fh_CivilLiberty, 4: fh_PoliticalRight, 5: fh_RuleOfLaw,
## #   6: fh_Status
```

```
names(d_gdp_edu_fh)
```

```
##  [1] "country_text_id"   "year"              "gdp"
##  [4] "gdppc"             "edu_15"            "edu_gini"
##  [7] "fh_CivilLiberty"   "fh_PoliticalRight" "fh_RuleOfLaw"
## [10] "fh_Status"
```

These are error-prone operations

- Do `bind_rows` and `bind_cols` ONLY WHEN you know for sure that there will not be a mismatch!
- If you have any slightest doubt, don't use them.

3. Join Tables

- `left_join`: Merge and only keep observations whose identifiers (matching keys) appear in the left-hand-side table.

- `right_join`: Merge and only keep observations whose identifiers (matching keys) appear in the right-hand-side table.
- `inner_join`: Merge and only keep observations whose identifiers (matching keys) appear in both tables.
- `full_join`: Merge and keep observations whose identifiers (matching keys) appear either table.
- `anti_join`: Filter out observations whose identifiers (matching keys) appear in the right-hand-side table
- `semi_join`: Filter out observations whose identifiers (matching keys) do not appear in the right-hand-side table

Task 1: The Case

Join two datasets from the V-Dem data using the above different `join_` functions

- *GDP* data from **2000-2022**
- *GDP per capita* data from **1984 to 2010**

```
# Setup
d_gdp_2000_2022 <- d |> filter(year %in% 2000:2022) |>
  select(country_text_id, year, e_gdp) |> rename("gdp" = "e_gdp")

d_gdppc_1984_2010 <- d |> filter(year %in% 1984:2010) |>
  select(country_text_id, year, e_gdppc) |> rename("gdppc" = "e_gdppc")

d_gdp_2000_2022 |> print(n = 2)
```

```
## # A tibble: 4,099 x 3
##   country_text_id  year     gdp
##   <chr>           <dbl>   <dbl>
## 1 MEX              2000 145206.
## 2 MEX              2001 146993.
## # ... with 4,097 more rows
```

```
d_gdppc_1984_2010 |> print(n = 2)
```

```
## # A tibble: 4,641 x 3
##   country_text_id  year gdppc
##   <chr>           <dbl> <dbl>
## 1 MEX              1984  11.7
## 2 MEX              1985  11.5
## # ... with 4,639 more rows
```

left_join

```
d_lj <- d_gdp_2000_2022 |>
  left_join(d_gdppc_1984_2010, by = c("country_text_id", "year"))

d_lj |> print(n = 2)
```

```
## # A tibble: 4,099 x 4
##   country_text_id  year     gdp gdppc
##   <chr>           <dbl>   <dbl> <dbl>
## 1 MEX              2000 145206.  13.7
## 2 MEX              2001 146993.  13.6
## # ... with 4,097 more rows
```

```
unique(d_lj$year) |> sort()
```

```
##  [1] 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010 2011 2012 2013 2014
## [16] 2015 2016 2017 2018 2019 2020 2021 2022
```

right_join

```
d_rj <- d_gdp_2000_2022 |>
  right_join(d_gdppc_1984_2010, by = c("country_text_id", "year"))

d_rj |> print(n = 2)
```

```
## # A tibble: 4,641 x 4
##   country_text_id  year     gdp gdppc
##   <chr>           <dbl>   <dbl> <dbl>
## 1 MEX              2000 145206.  13.7
## 2 MEX              2001 146993.  13.6
## # ... with 4,639 more rows
```

```
unique(d_rj$year) |> sort()
```

```
##  [1] 1984 1985 1986 1987 1988 1989 1990 1991 1992 1993 1994 1995 1996 1997 1998
## [16] 1999 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010
```

inner_join

```
d_ij <- d_gdp_2000_2022 |>
  inner_join(d_gdppc_1984_2010, by = c("country_text_id", "year"))

d_ij |> print(n = 2)
```

```
## # A tibble: 1,951 x 4
##   country_text_id  year     gdp gdppc
##   <chr>           <dbl>   <dbl> <dbl>
## 1 MEX              2000 145206.  13.7
## 2 MEX              2001 146993.  13.6
## # ... with 1,949 more rows
```

```r
unique(d_ij$year) |> sort()
```

```
##  [1] 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010
```

full_join

```r
d_fj <- d_gdp_2000_2022 |>
  full_join(d_gdppc_1984_2010, by = c("country_text_id", "year"))

d_fj |> print(n = 2)
```

```
## # A tibble: 6,789 x 4
##   country_text_id  year     gdp gdppc
##   <chr>           <dbl>   <dbl> <dbl>
## 1 MEX              2000 145206.  13.7
## 2 MEX              2001 146993.  13.6
## # ... with 6,787 more rows
```

```r
unique(d_fj$year) |> sort()
```

```
##  [1] 1984 1985 1986 1987 1988 1989 1990 1991 1992 1993 1994 1995 1996 1997 1998
## [16] 1999 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010 2011 2012 2013
## [31] 2014 2015 2016 2017 2018 2019 2020 2021 2022
```

semi_join

```r
d_sj <- d_gdp_2000_2022 |>
  semi_join(d_gdppc_1984_2010, by = c("country_text_id", "year"))

d_sj |> print(n = 2)
```

```
## # A tibble: 1,951 x 3
##   country_text_id  year     gdp
##   <chr>           <dbl>   <dbl>
## 1 MEX              2000 145206.
## 2 MEX              2001 146993.
## # ... with 1,949 more rows
```

```r
unique(d_sj$year) |> sort()
```

```
##  [1] 2000 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010
```

anti_join

```r
d_aj <- d_gdp_2000_2022 |>
  anti_join(d_gdppc_1984_2010, by = c("country_text_id", "year"))

d_aj |> print(n = 2)
```

```
## # A tibble: 2,148 x 3
##   country_text_id  year     gdp
##   <chr>           <dbl>   <dbl>
## 1 MEX              2011 185824.
## 2 MEX              2012 192272.
## # ... with 2,146 more rows
```

```r
unique(d_aj$year) |> sort()
```

```
##  [1] 2011 2012 2013 2014 2015 2016 2017 2018 2019 2020 2021 2022
```

3. Join by Identifiers with Different Variable Names

two options: (1) Rename it beforehand, (2) specify the `by =` argument differently.

```r
# I make an artificial example whose variable name of a matching
# identifier is different from d_gdp_2020_2022.
d_gdppc_1984_2010_t <- d_gdppc_1984_2010 |>
  rename("country_id" = "country_text_id")
```

```r
# Option 1: Rename the variables beforehand
d_aj_t <- d_gdp_2000_2022 |>
  rename("country_id" = "country_text_id") |>
  anti_join(d_gdppc_1984_2010_t, by = c("country_id", "year"))

# Option 2: Specify the "by =" argument with a *named vector*
d_aj_t_2 <- d_gdp_2000_2022 |>
  anti_join(d_gdppc_1984_2010_t,
            by = c("country_text_id" = "country_id",
                   "year" = "year"))
```

4. Many-to-One Join: Repeat!

Calculate each country's average 1984-2010 *GDP per capita* and merge it with our annual GDP data from 2000 to 2022.

```r
d_gdppc_1984_2010_avg <- d_gdppc_1984_2010 |> group_by(country_text_id) |>
  summarise(gdppc_1984to2010 = mean(gdppc, na.rm = TRUE))
d_gdppc_1984_2010_avg |> print(n = 2)
```

```
## # A tibble: 180 x 2
##   country_text_id gdppc_1984to2010
##   <chr>                      <dbl>
## 1 AFG                         1.22
## 2 AGO                         3.35
## # ... with 178 more rows
```

```r
d_lj_ManyToOne <- d_gdp_2000_2022 |>
  left_join(d_gdppc_1984_2010_avg, by = "country_text_id")
d_lj_ManyToOne |> print(n = 2)
```

```
## # A tibble: 4,099 x 4
##   country_text_id  year     gdp gdppc_1984to2010
##   <chr>           <dbl>   <dbl>            <dbl>
## 1 MEX              2000 145206.             12.8
## 2 MEX              2001 146993.             12.8
## # ... with 4,097 more rows
```

5. Save Outputs

- `.csv` "comma-separated values," readable by Excel or a text editor
- `.rds` "R data serialization," readable by R only

```r
# Save to a .csv file
write_csv(d_gdp_1945_2022, "Lecture_06/data/gdp_1945_2002.csv")

# Save to a .rds file
saveRDS(d_gdp_1945_2022, "Lecture_06/data/gdp_1945_2002.rds")
```

Saving Your Outputs after Data Wrangling

```r
# Read a .csv file
d_read_1 <- read_csv("Lecture_06/data/gdp_1945_2002.csv")

# Read a .rds file
d_read_2 <- readRDS("Lecture_06/data/gdp_1945_2002.rds")
```